

Computer Arithmetic

Professor: Yang Peng

The slides are re-produced by the courtesy of
Dr. Arnie Berger and Dr. Wooyoung Kim

Topic

Computer Arithmetic

- IEEE floating point
- Chapter 2.4, 2.5 (Null)

How to represent a real number
in binary?

From Real to Binary Numbers

- Let's convert decimal number **3.8125** to a binary number
 - Integer** part: the same as the integer binary $11_2 = 3$
 - Fractional** part:
 - Multiply the fraction by two**
 - Write down the **integer part on right**
 - Repeat 1 and 2 **until there is no fractional part on left**
 - Read the integer part on right, from top to bottom**

$$0.8125 * 2 = 0.625 + 1$$

$$0.625 * 2 = 0.25 + 1$$

$$0.25 * 2 = 0.5 + 0$$

$$0.5 * 2 = 0.0 + 1$$

0.0 STOP HERE

$$3.8125_{10} = 0011.1101_2$$

From Binary to Real Numbers

Binary $I_m I_{m-1} \dots I_1 I_0 \cdot F_1 F_2 F_3 \dots F_{n-1} F_n =$

Decimal $I \cdot 2^m + I \cdot 2^{m-1} + \dots + I \cdot 2^0 + F \cdot 2^{-1} + F \cdot 2^{-2}$
 $+ \dots + F \cdot 2^{-(n-1)} + F \cdot 2^{-n}$

$$10.101_2 = 1 \cdot 2^1 + 0 \cdot 2^0 + 1 \cdot 2^{-1} + 0 \cdot 2^{-2} + 1 \cdot 2^{-3}$$

$$= 2 + 0.5 + 0.125 = 2.625$$

Real Numbers and Errors

- Many fractions are **repeating infinitely**

E.g., convert 0.6 to binary

Integer

$$0.6 * 2 = 0.2 \text{ -----}1$$

$$0.2 * 2 = 0.4 \text{ -----}0$$

$$0.4 * 2 = 0.8 \text{ -----}0$$

$$0.8 * 2 = 0.6 \text{ -----}1$$

$$0.6 * 2 = 0.2 \text{ -----}1$$

$$0.2 * 2 = 0.4 \text{ -----}0$$

$$0.4 * 2 = 0.8 \text{ -----}0$$

$$0.8 * 2 = 0.6 \text{ -----}1$$

So, $0.6 \rightarrow 0.1001100110011001.....$ (will be repeated infinitely)

Rounding and Truncation

- Keep the number of bits **finite**
 - **Truncation**: The simplest technique – just drop unwanted bits
E.g., $0.1101101 \rightarrow 0.1101$
 - **Rounding**: Better technique, but a bit complicated

If the **value of the lost digits** is **greater than half of the least-significant bit of the retained digits**, add 1 to the LSB; otherwise drop.

E.g., 0.1101101 : If I want to lose the last three bits, what shall I do?

$$0.1101101 = 0.1101 + 0.0000101$$

$$\rightarrow 0.1101 + 0.0001$$

$$= 0.11\mathbf{10}$$

LSB of retained digits: $0.0001 = 2^{-4}$

Lost digits: $0.0000101 = (2^{-5} + 2^{-7}) > 2^{-4} / 2$

How to represent a real number
in a computer system?

Real Numbers in a Computer System

- Two main approaches: Fixed-point vs. Floating-point
- Fixed-point representation (**NOT used** now)
 - Divide the bits into **integer** part and **fraction** part
 - The “Point” is fixed
 - Easier but less flexible
- Floating-point representation (**IEEE standard**)
 - Divide the bits into **sign**, **exponent** and **mantissa**
 - The “Point” is floating
 - Match with scientific notation
 - Flexible but more complex

Fixed-Point Representation

- Fixed-point representation
 - Divide the bits for integer part and fraction part
 - For example, $3.625_{10} = 11.101_2$



Integer part

Fractional part

- Not flexible
 - What if you really need to represent $1.984 * 10^{(-123)}$ in computer?
 - How many bits will be needed? (more than 372 bits)

Floating-Point Representation

- Floating-point representation
 - Divide the bits into **sign**, **exponent** and **mantissa**
- IEEE floating-point format
 1. IEEE **short real** or **single precision**: **32** bits



2. IEEE **long real** or **double precision**: **64** bits



Floating-Point Representation

- Single Precision

Steps to convert a real number to IEEE **Single Precision** floating-point representation

1. Convert decimal to binary
2. Normalize: moving the point left or right
3. Add 127 to the exponent
4. Mantissa is the one **after the floating point** in the normalized form
 - If the mantissa part is less than 23 bits, add zeros at the end
5. Put the corresponding numbers into each field

IEEE Floating point converter:

<http://babbage.cs.qc.cuny.edu/IEEE-754.old/Decimal.html>

Floating-Point Representation

- Single Precision

- 32-bit (single precision) format

Sign (1)	Exponent (8)	Mantissa(23)
----------	--------------	--------------

- Let's represent a real number to floating-point format

E.g., -3.8125_{10}

$= -11.1101_2$ (note that the integer part is **not** 2's complement)

$= -1.11101 \cdot 2^1$ (**normalize**: scientific notation)

- **Sign** bit = 1, because this is a negative number
- **Exponent** bits = 1 + 127 (biased) = 128 = 10000000_2
- **Mantissa** bits: 111 0100 0000 0000 0000 0000

Therefore, the real number in floating-point representation is:

$1 \ 100 \ 0000 \ 0 \ 111 \ 01 \ 00 \ 0000 \ 0000 \ 0000 \ 0000_2 = \$C0740000$

Floating-Point Representation

- Single Precision

- **Normalization**
 - “1” shall always appear as an integer part
 - No need to represent this bit in the format -> save one bit
- **Biased exponent**
 - The exponent has 8 bits, meaning it can range from -127 to 127 (Here we assume that -128 will never appear).
 - Therefore, if we add 127 to the exponent, it will always be a non-negative number.
 - Assuming such a representation, 0~254 is then available for the exponent field. How about 255?

Floating-Point Representation

- Double Precision

Steps to convert a real number to IEEE **Double Precision** floating-point representation

1. Convert decimal to binary
2. Normalize: moving the point left or right
3. Add **1023** to the exponent
4. Mantissa is the one **after the floating point** in the normalized form
 - If the mantissa part is less than 52 bits, add zeros at the end
5. Put the corresponding numbers into each field

IEEE Floating point converter:

<http://babbage.cs.qc.cuny.edu/IEEE-754.old/Decimal.html>

Floating-Point Representation

- Double Precision

- **64-bit (double precision) format**

Sign (1)	Exponent (11)	Mantissa(52)
----------	---------------	--------------

- Let's represent a real number to floating-point format

E.g., -3.8125_{10}

$= -11.1101_2$ (note that the integer part is **not** 2's complement)

$= -1.11101 \cdot 2^1$ (normalize: scientific notation)

- **Sign** bit = 1, since negative
- **Exponent** = $1 + 1023$ (biased) = $1024 = 100\ 0000\ 0000_2$
- **Mantissa**: $1110\ 1000\ 0000\ 0000\ 0000\ 0000\ 0000\ 0000\ 0000\ 0000\ 0000\ 0000\ 0000\ 0000\ 0000\ 0000$

Therefore in floating-point representation,

$1\ 100\ 0000\ 0000\ 1110\ 1000\ 0000\ 0000\ 0000\ 0000\ 0000\ 0000\ 0000\ 0000\ 0000\ 0000_2$
 $= \$C00E800000000000$

More about real numbers

- Why using biased exponent?
 - **Effect**: changing **negative** exponent value **to positive** value
 - **Motivation**: for quick comparison (bit-by-bit) of two real numbers
- Why adding 127 for single-precision floating numbers?
 - **Effect**: positive numbers in the range of **0 to 254**
 - **Motivation**: reserve 255 for **special number usage**

Sign	Exponent (e)	Fraction (f)	Value
0	00...00	00...00	+0
0	00...00	00...01 ⋮ 11...11	Positive Denormalized Real $0.f \times 2^{(-b+1)}$
0	00...01 ⋮ 11...10	XX...XX	Positive Normalized Real $1.f \times 2^{(e-b)}$
0	11...11	00...00	$+\infty$
0	11...11	00...01 ⋮ 01...11	SNaN
0	11...11	1X...XX	QNaN
1	00...00	00...00	-0
1	00...00	00...01 ⋮ 11...11	Negative Denormalized Real $-0.f \times 2^{(-b+1)}$
1	00...01 ⋮ 11...10	XX...XX	Negative Normalized Real $-1.f \times 2^{(e-b)}$
1	11...11	00...00	$-\infty$
1	11...11	00...01 ⋮ 01...11	SNaN
1	11...11	1X...XX	QNaN

- **NaN: Not A Number**
- **QNaN: Quiet NaN**
 - generated from an operation when the result is not mathematically defined
 - denote ***indeterminate*** operations
- **SNaN: Signaling NaN**
 - used to signal an exception when used in operations
 - can be to assign to uninitialized variables to trap premature usage
 - denote ***invalid*** operations